

01_dclt_image_classification

April 28, 2025

M9 · Deep Learning aplicada - Visión artificial (Intel Scenes)

Clasificación multiclase 150×150 · 6 etiquetas

Proyecto desarrollado por Diego Cesar Lerma Torres para IMF Smart Education

Caso práctico del módulo **M9**

Deep learning aplicada: NLP y visión artificial

Del Master en Inteligencia Artificial

Este proyecto aplica técnicas avanzadas de **deep learning** para clasificar imágenes de escenas, combinando diversas áreas del conocimiento en inteligencia artificial.

1 Dataset

Intel Image Classification Dataset

Este conjunto de datos contiene imágenes de escenas naturales alrededor del mundo, capturadas en distintas condiciones y escenarios. El objetivo es clasificar cada imagen en una de seis categorías diferentes.

- **Número de imágenes:** Aproximadamente 25,000.
- **Dimensiones:** 150x150 píxeles por imagen.
- **Categorías:**
 - 0: Buildings (edificios)
 - 1: Forest (bosque)
 - 2: Glacier (glaciar)
 - 3: Mountain (montaña)
 - 4: Sea (mar)
 - 5: Street (calle)
- **Estructura del dataset:**
 - **Entrenamiento (Train):** ~14,000 imágenes.
 - **Pruebas (Test):** ~3,000 imágenes.
 - **Predicción (Prediction):** ~7,000 imágenes.

Cada partición está disponible en archivos comprimidos independientes.

Fuente: El dataset fue inicialmente publicado en [Analytics Vidhya](#) por Intel como parte de un desafío de clasificación de imágenes.

Objetivo del proyecto:

El propósito principal de este conjunto de datos es servir de base para el entrenamiento de redes neuronales profundas (CNNs) capaces de clasificar escenas naturales con alta precisión, fortaleciendo habilidades en visión artificial y aprendizaje profundo.

Agradecimientos a Intel y Analytics Vidhya por proporcionar este valioso recurso para la comunidad.

2 Configuración inicial

```
[3]: # !pip install -r requirements.txt
```

```
# En local  
%pip install -r requirements.txt
```

```
Requirement already satisfied: pandas in ./venv/lib/python3.12/site-packages  
(from -r requirements.txt (line 1)) (2.2.3)  
Requirement already satisfied: numpy in ./venv/lib/python3.12/site-packages  
(from -r requirements.txt (line 2)) (2.1.3)  
Requirement already satisfied: matplotlib in ./venv/lib/python3.12/site-  
packages (from -r requirements.txt (line 3)) (3.10.1)  
Requirement already satisfied: scikit-learn in ./venv/lib/python3.12/site-  
packages (from -r requirements.txt (line 4)) (1.6.1)  
Requirement already satisfied: keras-tuner in ./venv/lib/python3.12/site-  
packages (from -r requirements.txt (line 5)) (1.4.7)  
Requirement already satisfied: tensorflow[and-cuda] in  
./venv/lib/python3.12/site-packages (from -r requirements.txt (line 6))  
(2.19.0)  
Requirement already satisfied: python-dateutil>=2.8.2 in  
./venv/lib/python3.12/site-packages (from pandas->-r requirements.txt (line 1))  
(2.9.0.post0)  
Requirement already satisfied: pytz>=2020.1 in ./venv/lib/python3.12/site-  
packages (from pandas->-r requirements.txt (line 1)) (2025.2)  
Requirement already satisfied: tzdata>=2022.7 in ./venv/lib/python3.12/site-  
packages (from pandas->-r requirements.txt (line 1)) (2025.2)  
Requirement already satisfied: contourpy>=1.0.1 in ./venv/lib/python3.12/site-  
packages (from matplotlib->-r requirements.txt (line 3)) (1.3.2)  
Requirement already satisfied: cycycler>=0.10 in ./venv/lib/python3.12/site-  
packages (from matplotlib->-r requirements.txt (line 3)) (0.12.1)  
Requirement already satisfied: fonttools>=4.22.0 in ./venv/lib/python3.12/site-  
packages (from matplotlib->-r requirements.txt (line 3)) (4.57.0)  
Requirement already satisfied: kiwisolver>=1.3.1 in ./venv/lib/python3.12/site-  
packages (from matplotlib->-r requirements.txt (line 3)) (1.4.8)  
Requirement already satisfied: packaging>=20.0 in ./venv/lib/python3.12/site-  
packages (from matplotlib->-r requirements.txt (line 3)) (25.0)  
Requirement already satisfied: pillow>=8 in ./venv/lib/python3.12/site-packages  
(from matplotlib->-r requirements.txt (line 3)) (11.2.1)  
Requirement already satisfied: pyparsing>=2.3.1 in ./venv/lib/python3.12/site-
```

packages (from matplotlib->-r requirements.txt (line 3)) (3.2.3)
 Requirement already satisfied: scipy>=1.6.0 in ./venv/lib/python3.12/site-packages (from scikit-learn->-r requirements.txt (line 4)) (1.15.2)
 Requirement already satisfied: joblib>=1.2.0 in ./venv/lib/python3.12/site-packages (from scikit-learn->-r requirements.txt (line 4)) (1.4.2)
 Requirement already satisfied: threadpoolctl>=3.1.0 in ./venv/lib/python3.12/site-packages (from scikit-learn->-r requirements.txt (line 4)) (3.6.0)
 Requirement already satisfied: keras in ./venv/lib/python3.12/site-packages (from keras-tuner->-r requirements.txt (line 5)) (3.9.2)
 Requirement already satisfied: requests in ./venv/lib/python3.12/site-packages (from keras-tuner->-r requirements.txt (line 5)) (2.32.3)
 Requirement already satisfied: kt-legacy in ./venv/lib/python3.12/site-packages (from keras-tuner->-r requirements.txt (line 5)) (1.0.5)
 Requirement already satisfied: absl-py>=1.0.0 in ./venv/lib/python3.12/site-packages (from tensorflow[and-cuda]->-r requirements.txt (line 6)) (2.2.2)
 Requirement already satisfied: astunparse>=1.6.0 in ./venv/lib/python3.12/site-packages (from tensorflow[and-cuda]->-r requirements.txt (line 6)) (1.6.3)
 Requirement already satisfied: flatbuffers>=24.3.25 in ./venv/lib/python3.12/site-packages (from tensorflow[and-cuda]->-r requirements.txt (line 6)) (25.2.10)
 Requirement already satisfied: gast!=0.5.0,!0.5.1,!0.5.2,>=0.2.1 in ./venv/lib/python3.12/site-packages (from tensorflow[and-cuda]->-r requirements.txt (line 6)) (0.6.0)
 Requirement already satisfied: google-pasta>=0.1.1 in ./venv/lib/python3.12/site-packages (from tensorflow[and-cuda]->-r requirements.txt (line 6)) (0.2.0)
 Requirement already satisfied: libclang>=13.0.0 in ./venv/lib/python3.12/site-packages (from tensorflow[and-cuda]->-r requirements.txt (line 6)) (18.1.1)
 Requirement already satisfied: opt-einsum>=2.3.2 in ./venv/lib/python3.12/site-packages (from tensorflow[and-cuda]->-r requirements.txt (line 6)) (3.4.0)
 Requirement already satisfied:
 protobuf!=4.21.0,!4.21.1,!4.21.2,!4.21.3,!4.21.4,!4.21.5,<6.0.0dev,>=3.20.3 in ./venv/lib/python3.12/site-packages (from tensorflow[and-cuda]->-r requirements.txt (line 6)) (5.29.4)
 Requirement already satisfied: setuptools in ./venv/lib/python3.12/site-packages (from tensorflow[and-cuda]->-r requirements.txt (line 6)) (80.0.0)
 Requirement already satisfied: six>=1.12.0 in ./venv/lib/python3.12/site-packages (from tensorflow[and-cuda]->-r requirements.txt (line 6)) (1.17.0)
 Requirement already satisfied: termcolor>=1.1.0 in ./venv/lib/python3.12/site-packages (from tensorflow[and-cuda]->-r requirements.txt (line 6)) (3.0.1)
 Requirement already satisfied: typing-extensions>=3.6.6 in ./venv/lib/python3.12/site-packages (from tensorflow[and-cuda]->-r requirements.txt (line 6)) (4.13.2)
 Requirement already satisfied: wrapt>=1.11.0 in ./venv/lib/python3.12/site-packages (from tensorflow[and-cuda]->-r requirements.txt (line 6)) (1.17.2)
 Requirement already satisfied: grpcio<2.0,>=1.24.3 in ./venv/lib/python3.12/site-packages (from tensorflow[and-cuda]->-r

```

requirements.txt (line 6)) (1.71.0)
Requirement already satisfied: tensorboard~=2.19.0 in
./venv/lib/python3.12/site-packages (from tensorflow[and-cuda]->-r
requirements.txt (line 6)) (2.19.0)
Requirement already satisfied: h5py>=3.11.0 in ./venv/lib/python3.12/site-
packages (from tensorflow[and-cuda]->-r requirements.txt (line 6)) (3.13.0)
Requirement already satisfied: ml-dtypes<1.0.0,>=0.5.1 in
./venv/lib/python3.12/site-packages (from tensorflow[and-cuda]->-r
requirements.txt (line 6)) (0.5.1)
Requirement already satisfied: nvidia-cublas-cu12==12.5.3.2 in
./venv/lib/python3.12/site-packages (from tensorflow[and-cuda]->-r
requirements.txt (line 6)) (12.5.3.2)
Requirement already satisfied: nvidia-cuda-cupti-cu12==12.5.82 in
./venv/lib/python3.12/site-packages (from tensorflow[and-cuda]->-r
requirements.txt (line 6)) (12.5.82)
Requirement already satisfied: nvidia-cuda-nvcc-cu12==12.5.82 in
./venv/lib/python3.12/site-packages (from tensorflow[and-cuda]->-r
requirements.txt (line 6)) (12.5.82)
Requirement already satisfied: nvidia-cuda-nvrtc-cu12==12.5.82 in
./venv/lib/python3.12/site-packages (from tensorflow[and-cuda]->-r
requirements.txt (line 6)) (12.5.82)
Requirement already satisfied: nvidia-cuda-runtime-cu12==12.5.82 in
./venv/lib/python3.12/site-packages (from tensorflow[and-cuda]->-r
requirements.txt (line 6)) (12.5.82)
Requirement already satisfied: nvidia-cudnn-cu12==9.3.0.75 in
./venv/lib/python3.12/site-packages (from tensorflow[and-cuda]->-r
requirements.txt (line 6)) (9.3.0.75)
Requirement already satisfied: nvidia-cufft-cu12==11.2.3.61 in
./venv/lib/python3.12/site-packages (from tensorflow[and-cuda]->-r
requirements.txt (line 6)) (11.2.3.61)
Requirement already satisfied: nvidia-curand-cu12==10.3.6.82 in
./venv/lib/python3.12/site-packages (from tensorflow[and-cuda]->-r
requirements.txt (line 6)) (10.3.6.82)
Requirement already satisfied: nvidia-cusolver-cu12==11.6.3.83 in
./venv/lib/python3.12/site-packages (from tensorflow[and-cuda]->-r
requirements.txt (line 6)) (11.6.3.83)
Requirement already satisfied: nvidia-cusparse-cu12==12.5.1.3 in
./venv/lib/python3.12/site-packages (from tensorflow[and-cuda]->-r
requirements.txt (line 6)) (12.5.1.3)
Requirement already satisfied: nvidia-nccl-cu12==2.23.4 in
./venv/lib/python3.12/site-packages (from tensorflow[and-cuda]->-r
requirements.txt (line 6)) (2.23.4)
Requirement already satisfied: nvidia-nvjitlink-cu12==12.5.82 in
./venv/lib/python3.12/site-packages (from tensorflow[and-cuda]->-r
requirements.txt (line 6)) (12.5.82)
Requirement already satisfied: charset-normalizer<4,>=2 in
./venv/lib/python3.12/site-packages (from requests->keras-tuner->-r
requirements.txt (line 5)) (3.4.1)

```

Requirement already satisfied: idna<4,>=2.5 in ./venv/lib/python3.12/site-packages (from requests->keras-tuner->-r requirements.txt (line 5)) (3.10)

Requirement already satisfied: urllib3<3,>=1.21.1 in ./venv/lib/python3.12/site-packages (from requests->keras-tuner->-r requirements.txt (line 5)) (2.4.0)

Requirement already satisfied: certifi>=2017.4.17 in ./venv/lib/python3.12/site-packages (from requests->keras-tuner->-r requirements.txt (line 5)) (2025.4.26)

Requirement already satisfied: markdown>=2.6.8 in ./venv/lib/python3.12/site-packages (from tensorboard~=2.19.0->tensorflow[and-cuda]->-r requirements.txt (line 6)) (3.8)

Requirement already satisfied: tensorboard-data-server<0.8.0,>=0.7.0 in ./venv/lib/python3.12/site-packages (from tensorboard~=2.19.0->tensorflow[and-cuda]->-r requirements.txt (line 6)) (0.7.2)

Requirement already satisfied: werkzeug>=1.0.1 in ./venv/lib/python3.12/site-packages (from tensorboard~=2.19.0->tensorflow[and-cuda]->-r requirements.txt (line 6)) (3.1.3)

Requirement already satisfied: wheel<1.0,>=0.23.0 in ./venv/lib/python3.12/site-packages (from astunparse>=1.6.0->tensorflow[and-cuda]->-r requirements.txt (line 6)) (0.45.1)

Requirement already satisfied: rich in ./venv/lib/python3.12/site-packages (from keras->keras-tuner->-r requirements.txt (line 5)) (14.0.0)

Requirement already satisfied: namex in ./venv/lib/python3.12/site-packages (from keras->keras-tuner->-r requirements.txt (line 5)) (0.0.9)

Requirement already satisfied: optree in ./venv/lib/python3.12/site-packages (from keras->keras-tuner->-r requirements.txt (line 5)) (0.15.0)

Requirement already satisfied: MarkupSafe>=2.1.1 in ./venv/lib/python3.12/site-packages (from werkzeug>=1.0.1->tensorboard~=2.19.0->tensorflow[and-cuda]->-r requirements.txt (line 6)) (3.0.2)

Requirement already satisfied: markdown-it-py>=2.2.0 in ./venv/lib/python3.12/site-packages (from rich->keras->keras-tuner->-r requirements.txt (line 5)) (3.0.0)

Requirement already satisfied: pygments<3.0.0,>=2.13.0 in ./venv/lib/python3.12/site-packages (from rich->keras->keras-tuner->-r requirements.txt (line 5)) (2.19.1)

Requirement already satisfied: mdurl~=0.1 in ./venv/lib/python3.12/site-packages (from markdown-it-py>=2.2.0->rich->keras->keras-tuner->-r requirements.txt (line 5)) (0.1.2)

Note: you may need to restart the kernel to use updated packages.

```
[ ]: # Subir el kaggle.json, en caso de querer descargar directamente la base de
      ↪ datos de su origen

#from google.colab import files
#files.upload()
```

```
[3]: # Mover el archivo descargado a su lugar
```

```
#!/mkdir -p ~/.kaggle
#!/mv kaggle.json ~/.kaggle/
#!/chmod 600 ~/.kaggle/kaggle.json
```

```
[4]: import os
import json
import zipfile
import random
from pathlib import Path

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

import tensorflow as tf
from tensorflow.keras import layers, models, mixed_precision
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint, \
    ReduceLROnPlateau
from tensorflow.keras.preprocessing.image import ImageDataGenerator

import keras_tuner as kt

mixed_precision.set_global_policy('mixed_float16')

SEED = 42
IMG_SIZE = (150, 150)
BATCH_SIZE = 16
EPOCHS = 12
VAL_SPLIT = 0.20
TEST_SPLIT = 0.25
AUTOTUNE = tf.data.AUTOTUNE
PROJECT_ROOT = Path('.')
DATA_DIR = PROJECT_ROOT/'intel_scenes'
FULL_DATA = DATA_DIR / 'data_all'
MODELS_DIR = PROJECT_ROOT/'models'
HIST_DIR = PROJECT_ROOT/'history'
for d in (DATA_DIR, FULL_DATA, MODELS_DIR, HIST_DIR):
    d.mkdir(parents=True, exist_ok=True)

random.seed(SEED);
np.random.seed(SEED);
tf.random.set_seed(SEED)
print(tf.__version__)
```

2.19.0

Ejecutar solo la primera vez. Descargar y descomprimir desde Kaggle

Las siguientes veces, esta sección debe comentarse u omitirse su ejecución

```
[ ]: !kaggle datasets download -d puneet6060/intel-image-classification -p_
    ↪{PROJECT_ROOT}
```

```
[6]: with zipfile.ZipFile(PROJECT_ROOT/'intel-image-classification.zip') as z: z.
    ↪extractall(DATA_DIR)
```

3 1. Descarga y unificación de carpetas seg_train + seg_test

```
[4]: """# 1. Unificación de carpetas seg_train + seg_test"""

import shutil

def merge_folders_corrected(src_root: Path, dst_root: Path):
    """
    Fusiona las imágenes de las subcarpetas de clase encontradas dentro de
    src_root/seg_train/seg_train/<clase> y src_root/seg_test/seg_test/<clase>
    en dst_root/<clase>.
    Elimina las carpetas originales seg_train/ y seg_test/ de src_root después.
    """
    print(f"Iniciando fusión de carpetas desde {src_root} hacia {dst_root}")
    dst_root.mkdir(parents=True, exist_ok=True)

    split_names = ['seg_train', 'seg_test']

    for split in split_names:
        split_base_dir = src_root / split
        print(f"Procesando división: {split}")

        nested_split_dir = split_base_dir / split

        if nested_split_dir.exists() and nested_split_dir.is_dir():
            print(f" Encontrada estructura anidada en: {nested_split_dir}")
            source_class_dir_parent = nested_split_dir
        elif split_base_dir.exists() and split_base_dir.is_dir():
            contains_class_dirs = any(item.is_dir() for item in split_base_dir.
            ↪iterdir())
            if contains_class_dirs:
                print(f" Estructura anidada no encontrada, usando directorio_
            ↪base: {split_base_dir}")
                source_class_dir_parent = split_base_dir
            else:
                print(f" Directorio {split_base_dir} no contiene_
            ↪subdirectorios de clase esperados. Omitiendo.")
```

```

        continue
    else:
        print(f" Directorio base {split_base_dir} no encontrado o no es un_
↳directorio. Omitiendo.")
        continue

    class_count = 0
    image_count = 0
    for class_dir in source_class_dir_parent.iterdir():
        if class_dir.is_dir():
            class_name = class_dir.name
            target_class_dir = dst_root / class_name
            target_class_dir.mkdir(parents=True, exist_ok=True)
            class_count += 1

            files_in_class = list(class_dir.iterdir())
            print(f" Procesando clase '{class_name}'_
↳({len(files_in_class)} archivos)...")
            for img_path in files_in_class:
                if img_path.is_file():
                    dst_img_path = target_class_dir / img_path.name
                    try:
                        shutil.move(str(img_path), str(dst_img_path))
                        image_count += 1
                    except Exception as e:
                        print(f" Error moviendo {img_path.name} a_
↳{target_class_dir}: {e}")

            print(f" Procesadas {class_count} clases y {image_count} imágenes para_
↳la división '{split}'.")

    print("\nLimpiando carpetas originales...")
    for split in split_names:
        original_split_dir = src_root / split
        if original_split_dir.exists():
            try:
                shutil.rmtree(original_split_dir)
                print(f" Eliminada carpeta: {original_split_dir}")
            except OSError as e:
                print(f" Error eliminando {original_split_dir}: {e}")
        else:
            print(f" La carpeta {original_split_dir} no existe, no se necesita_
↳eliminar.")

    print(f"\nFusión completada. Datos unificados en: {dst_root}")

```



```
[ ]: if DATA_DIR.exists():
    merge_folders_corrected(DATA_DIR, FULL_DATA)
else:
    print(f"ERROR: El directorio de origen {DATA_DIR} no existe. No se puede_
    ↪ejecutar la fusión.")

if FULL_DATA.exists():
    print("\nContenido del directorio unificado (data_all):")
    class_names_merged = sorted([d.name for d in FULL_DATA.iterdir() if d.
    ↪is_dir()])
    print(f"Clases encontradas: {class_names_merged}")
    total_images = 0
    for class_dir in FULL_DATA.iterdir():
        if class_dir.is_dir():
            count = len(list(class_dir.glob('*.*')))
            print(f"- {class_dir.name}: {count} imágenes")
            total_images += count
    print(f"Total imágenes unificadas: {total_images}")
else:
    print(f"ERROR: El directorio destino {FULL_DATA} no se creó correctamente.")
```

4 2. Generadores de imágenes

```
[6]: def build_dataset(directory, img_size=IMG_SIZE, batch_size=BATCH_SIZE,
    test_split=TEST_SPLIT, seed=SEED):

    ds_full = tf.keras.utils.image_dataset_from_directory(
        directory,
        image_size=img_size,
        batch_size=batch_size,
        shuffle=True,
        seed=seed,
        label_mode='int')

    class_names = ds_full.class_names

    total      = ds_full.cardinality().numpy()
    test_count = int(total * test_split)

    test_ds    = ds_full.take(test_count)
    train_val  = ds_full.skip(test_count)

    val_count  = int(train_val.cardinality().numpy() * VAL_SPLIT)
    val_ds     = train_val.take(val_count)
    train_ds   = train_val.skip(val_count)
```

```

train_ds = train_ds.cache().prefetch(AUTOTUNE)
val_ds   = val_ds.cache().prefetch(AUTOTUNE)
test_ds  = test_ds.cache().prefetch(AUTOTUNE)

return train_ds, val_ds, test_ds, class_names

```

```

[7]: train_ds, val_ds, test_ds, class_names = build_dataset(FULL_DATA)
NUM_CLASSES = len(class_names)
print(class_names)

```

Found 17034 files belonging to 6 classes.

```
['buildings', 'forest', 'glacier', 'mountain', 'sea', 'street']
```

```

[8]: for cls in sorted((FULL_DATA).iterdir()):
    print(f"{cls.name:<10} {len(list(cls.glob('*'))):5d} imágenes")

# Debe mostrarse que existen varias imagenes repartidas en 6 carpetas. De lo
↪ contrario, revisar antes de comenzar el entrenamiento.

```

buildings	2628 imágenes
forest	2745 imágenes
glacier	2957 imágenes
mountain	3037 imágenes
sea	2784 imágenes
street	2883 imágenes

5 3. Callbacks comunes

```

[9]: def common_callbacks(name):
    return [
        EarlyStopping(patience=2, restore_best_weights=True),
        ModelCheckpoint(MODELS_DIR/f'{name}.keras', save_best_only=True),
        ReduceLROnPlateau(patience=1, factor=0.3, verbose=1)
    ]

```

```

[10]: def get_optimizer(lr=1e-3):
    """Crea un optimizador Adam envuelto en LossScaleOptimizer."""
    base_opt = tf.keras.optimizers.Adam(learning_rate=lr)
    return mixed_precision.LossScaleOptimizer(base_opt)

```

6 4. Modelo 1 — CNN Base

```
[11]: def cnn_base(input_shape=IMG_SIZE+(3,), num_classes=NUM_CLASSES):
```

```
    model = tf.keras.Sequential([
        tf.keras.Input(shape=input_shape),
        layers.Rescaling(1./255),
        layers.Conv2D(32, 3, activation='relu'),
        layers.MaxPooling2D(),
        layers.Flatten(),
        layers.Dense(64, activation='relu'),
        layers.Dense(num_classes, activation='softmax')
    ])
    model.compile(optimizer=get_optimizer(1e-3),
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])
    return model
```

```
[12]: base_model = cnn_base()
history_base = base_model.fit(
    train_ds, epochs=EPOCHS, validation_data=val_ds,
    callbacks=common_callbacks('base_model'))
json.dump(history_base.history, open(HIST_DIR/'base_model.json', 'w'))
```

Epoch 1/12

```
WARNING: All log messages before absl::InitializeLog() is called are written to
STDERR
I0000 00:00:1745863652.858889    21159 service.cc:152] XLA service 0x7f2dcc0045e0
initialized for platform CUDA (this does not guarantee that XLA will be used).
Devices:
I0000 00:00:1745863652.858935    21159 service.cc:160]   StreamExecutor device
(0): NVIDIA GeForce RTX 4090, Compute Capability 8.9
2025-04-28 12:07:32.901910: I
tensorflow/compiler/mlir/tensorflow/utils/dump_mlir_util.cc:269] disabling MLIR
crash reproducer, set env var `MLIR_CRASH_REPRODUCER_DIRECTORY` to enable.
I0000 00:00:1745863653.011512    21159 cuda_dnn.cc:529] Loaded cuDNN version
90300

 33/640                2s 4ms/step -
accuracy: 0.1948 - loss: 8.3145

I0000 00:00:1745863654.645041    21159 device_compiler.h:188] Compiled cluster
using XLA! This line is logged at most once for the lifetime of the process.

640/640                0s 6ms/step -
accuracy: 0.4314 - loss: 2.7119

2025-04-28 12:07:39.328600: I
external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas
```

warning : Registers are spilled to local memory in function
'gemm_fusion_dot_56', 344 bytes spill stores, 344 bytes spill loads

```
640/640          8s 9ms/step -  
accuracy: 0.4316 - loss: 2.7099 - val_accuracy: 0.6922 - val_loss: 0.8303 -  
learning_rate: 0.0010  
Epoch 2/12  
640/640          0s 3ms/step -  
accuracy: 0.7771 - loss: 0.6218  
Epoch 2: ReduceLROnPlateau reducing learning rate to 0.0003000000142492354.  
640/640          2s 3ms/step -  
accuracy: 0.7772 - loss: 0.6216 - val_accuracy: 0.6588 - val_loss: 1.0801 -  
learning_rate: 0.0010  
Epoch 3/12  
637/640          0s 2ms/step -  
accuracy: 0.9318 - loss: 0.2339  
Epoch 3: ReduceLROnPlateau reducing learning rate to 9.000000427477062e-05.  
640/640          2s 3ms/step -  
accuracy: 0.9319 - loss: 0.2335 - val_accuracy: 0.7268 - val_loss: 0.8828 -  
learning_rate: 3.0000e-04
```

7 5. Modelo 2 — CNN Avanzada (más profundidad + Dropout + BatchNorm)

```
[13]: def cnn_advanced(input_shape=IMG_SIZE+(3,), num_classes=NUM_CLASSES):  
    inputs = layers.Input(shape=input_shape)  
    x = layers.Rescaling(1./255)(inputs)  
  
    for filters in [32, 64, 128]:  
        x = layers.Conv2D(filters, 3, padding='same', activation='relu')(x)  
        x = layers.BatchNormalization()(x)  
        x = layers.Conv2D(filters, 3, padding='same', activation='relu')(x)  
        x = layers.MaxPooling2D()(x)  
        x = layers.Dropout(0.25)(x)  
  
    x = layers.Flatten()(x)  
    x = layers.Dense(256, activation='relu')(x)  
    x = layers.Dropout(0.5)(x)  
    outputs = layers.Dense(num_classes, activation='softmax')(x)  
  
    model = tf.keras.Model(inputs, outputs, name='cnn_advanced')  
    model.compile(optimizer=get_optimizer(1e-3),  
                  loss='sparse_categorical_crossentropy',  
                  metrics=['accuracy'])  
    return model
```

```
[14]: adv_model = cnn_advanced()
      history_adv = adv_model.fit(
          train_ds, epochs=EPOCHS, validation_data=val_ds,
          callbacks=common_callbacks('advance_model'))
      json.dump(history_adv.history, open(HIST_DIR/'advance_model.json','w'))
```

Epoch 1/12

640/640 0s 11ms/step -
accuracy: 0.1826 - loss: 12.6964

2025-04-28 12:07:57.815448: I

external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas
warning : Registers are spilled to local memory in function
'gemm_fusion_dot_175', 4 bytes spill stores, 4 bytes spill loads

2025-04-28 12:07:57.881220: I

external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas
warning : Registers are spilled to local memory in function
'gemm_fusion_dot_168', 240 bytes spill stores, 240 bytes spill loads

640/640 16s 16ms/step -
accuracy: 0.1826 - loss: 12.6970 - val_accuracy: 0.1741 - val_loss: 13.3114 -
learning_rate: 0.0010

Epoch 2/12

631/640 0s 4ms/step -
accuracy: 0.1772 - loss: 13.2620

Epoch 2: ReduceLROnPlateau reducing learning rate to 0.0003000000142492354.

640/640 3s 5ms/step -
accuracy: 0.1771 - loss: 13.2622 - val_accuracy: 0.1741 - val_loss: nan -
learning_rate: 0.0010

Epoch 3/12

639/640 0s 4ms/step -
accuracy: 0.1746 - loss: 13.3041

Epoch 3: ReduceLROnPlateau reducing learning rate to 9.000000427477062e-05.

640/640 3s 5ms/step -
accuracy: 0.1746 - loss: 13.3043 - val_accuracy: 0.1741 - val_loss: 13.3114 -
learning_rate: 3.0000e-04

8 6. Modelo 3 — Hyperparameter Tuning (Keras Tuner)

```
[16]: def model_builder(hp):
      hp_filters = hp.Choice('filters', values=[32, 48, 64])
      hp_dense = hp.Int('dense_units', min_value=128, max_value=512, step=128)
      hp_dropout = hp.Float('dropout', 0.2, 0.5, step=0.1)
      hp_lr = hp.Choice('lr', [1e-2, 1e-3, 1e-4])
```

```

inputs = layers.Input(shape=IMG_SIZE + (3,))
x = layers.Rescaling(1./255)(inputs)

for filters in [hp_filters, hp_filters*2, hp_filters*4]:
    x = layers.Conv2D(filters, 3, padding='same', activation='relu')(x)
    x = layers.Conv2D(filters, 3, padding='same', activation='relu')(x)
    x = layers.MaxPooling2D()(x)
    x = layers.Dropout(hp_dropout)(x)

x = layers.Flatten()(x)
x = layers.Dense(hp_dense, activation='relu')(x)
x = layers.Dropout(hp_dropout)(x)
outputs = layers.Dense(NUM_CLASSES, activation='softmax')(x)

model = tf.keras.Model(inputs, outputs)

opt = get_optimizer(hp_lr)

model.compile(
    optimizer=opt,
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy'])
return model

```

```

[17]: tuner = kt.RandomSearch(
    model_builder,
    objective='val_accuracy',
    max_trials=10,
    executions_per_trial=1,
    directory='kt_logs',
    project_name='intel_hp',
    overwrite=True)

SEARCH_EPOCHS = 5

print(f"Iniciando búsqueda de hiperparámetros (max_trials={tuner.oracle.
↳max_trials}, epochs_per_trial={SEARCH_EPOCHS})...")

tuner.search(train_ds,
             epochs=SEARCH_EPOCHS,
             validation_data=val_ds,
             callbacks=[EarlyStopping(monitor='val_accuracy', patience=2)])

print("\nBúsqueda completada. Obteniendo y re-entrenando el mejor modelo...")

best_hp = tuner.get_best_hyperparameters(1)[0]
print("Mejores Hiperparámetros encontrados:")

```

```

print(best_hp.values)

best_hp_model = tuner.hypermodel.build(best_hp)

best_hp_model.save(MODELS_DIR / 'hp_best_structure_untrained.keras',
    ↪include_optimizer=False)

history_hp = best_hp_model.fit(
    train_ds,
    epochs=EPOCHS,
    validation_data=val_ds,
    callbacks=common_callbacks('hp')
)

json.dump(history_hp.history, open(HIST_DIR / 'hp.json', 'w'))
best_hp_model.save(MODELS_DIR / 'hp.keras')

```

Trial 10 Complete [00h 00m 20s]
val_accuracy: 0.18553459644317627

Best val_accuracy So Far: 0.8164308071136475
Total elapsed time: 00h 06m 31s

Búsqueda completada. Obteniendo y re-entrenando el mejor modelo...

Mejores Hiperparámetros encontrados:

{'filters': 64, 'dense_units': 384, 'dropout': 0.4, 'lr': 0.0001}

Epoch 1/12

640/640 18s 23ms/step -

accuracy: 0.4183 - loss: 1.3841 - val_accuracy: 0.6694 - val_loss: 0.8670 -
learning_rate: 1.0000e-04

Epoch 2/12

640/640 11s 17ms/step -

accuracy: 0.6420 - loss: 0.8925 - val_accuracy: 0.7559 - val_loss: 0.6730 -
learning_rate: 1.0000e-04

Epoch 3/12

640/640 12s 19ms/step -

accuracy: 0.7483 - loss: 0.6798 - val_accuracy: 0.8062 - val_loss: 0.5517 -
learning_rate: 1.0000e-04

Epoch 4/12

640/640 11s 17ms/step -

accuracy: 0.8059 - loss: 0.5513 - val_accuracy: 0.8156 - val_loss: 0.5107 -
learning_rate: 1.0000e-04

Epoch 5/12

640/640 11s 17ms/step -

accuracy: 0.8390 - loss: 0.4535 - val_accuracy: 0.8192 - val_loss: 0.4918 -
learning_rate: 1.0000e-04

Epoch 6/12

```

639/640          0s 16ms/step -
accuracy: 0.8657 - loss: 0.3922
Epoch 6: ReduceLROnPlateau reducing learning rate to 2.9999999242136255e-05.
640/640          11s 17ms/step -
accuracy: 0.8657 - loss: 0.3921 - val_accuracy: 0.8227 - val_loss: 0.5062 -
learning_rate: 1.0000e-04
Epoch 7/12
638/640          0s 15ms/step -
accuracy: 0.8914 - loss: 0.3031
Epoch 7: ReduceLROnPlateau reducing learning rate to 8.999999772640877e-06.
640/640          10s 16ms/step -
accuracy: 0.8914 - loss: 0.3030 - val_accuracy: 0.8322 - val_loss: 0.4964 -
learning_rate: 3.0000e-05

```

9 7. Modelo 4 — Transfer Learning + Fine Tuning

```

[18]: def transfer_finetune(base='MobileNetV2',
                             img_size=IMG_SIZE,
                             num_classes=NUM_CLASSES,
                             unfreeze_from=100,
                             hub_size=224):

    """Feature-extraction + fine-tuning con red pre-entrenada."""
    base_model = get_attr(tf.keras.applications, base)(
        include_top=False,
        weights='imagenet',
        input_shape=(hub_size, hub_size, 3)
    )
    base_model.trainable = False

    inputs = layers.Input(shape=img_size + (3,))
    x = layers.Resizing(hub_size, hub_size)(inputs)
    x = layers.Rescaling(1./255)(x)
    x = base_model(x, training=False)
    x = layers.GlobalAveragePooling2D()(x)
    x = layers.Dense(128, activation='relu')(x)
    outputs = layers.Dense(num_classes, activation='softmax')(x)

    model = tf.keras.Model(inputs, outputs, name=f'{base}_finetune')

    opt_fe = get_optimizer(1e-3)
    model.compile(optimizer=opt_fe,
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])
    hist_fe = model.fit(train_ds,
                        epochs=5,
                        validation_data=val_ds)

```



```

opt_ft = get_optimizer(1e-5)

base_model.trainable = True
print(f"Fine-tuning: Descongelando desde la capa {unfreeze_from}")
for i, layer in enumerate(base_model.layers):
    if i < unfreeze_from:
        layer.trainable = False
    else:
        if i % 10 == 0 or i >= len(base_model.layers) - 5:
            print(f" - Capa {i} ({layer.name}): Trainable = {layer.
↵trainable}")

print("\nRe-compilando modelo para fine-tuning con LR bajo...")
model.compile(optimizer=opt_ft,
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

print("Iniciando fase de fine-tuning...")
initial_epoch_ft = hist_fe.epoch[-1] + 1

hist_ft = model.fit(train_ds,
                   epochs=EPOCHS,
                   validation_data=val_ds,
                   callbacks=common_callbacks('fine_tuning'))

history = {k: hist_fe.history.get(k, []) + hist_ft.history[k]
          for k in hist_ft.history.keys()}
}
return model, history

```

```

[19]: tl_model, history_ft = transfer_finetune()
      json.dump(history_ft, open(HIST_DIR / 'fine_tuning.json', 'w'))
      tl_model.save(MODELS_DIR / 'fine_tuning.keras', include_optimizer=False)

```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/mobilenet_v2/mobilenet_v2_weights_tf_dim_ordering_tf_kernels_1.0_224_no_top.h5

9406464/9406464

0s

0us/step

Epoch 1/5

2025-04-28 12:16:07.045049: I

external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas

warning : Registers are spilled to local memory in function

'gemm_fusion_dot_1449', 4 bytes spill stores, 4 bytes spill loads

2025-04-28 12:16:07.148869: I
external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas
warning : Registers are spilled to local memory in function
'gemm_fusion_dot_1449_0', 200 bytes spill stores, 536 bytes spill loads

2025-04-28 12:16:07.220225: I
external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas
warning : Registers are spilled to local memory in function
'gemm_fusion_dot_1449', 116 bytes spill stores, 116 bytes spill loads

2025-04-28 12:16:07.481231: I
external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas
warning : Registers are spilled to local memory in function
'gemm_fusion_dot_1449', 5616 bytes spill stores, 5612 bytes spill loads

2025-04-28 12:16:07.614588: I
external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas
warning : Registers are spilled to local memory in function
'gemm_fusion_dot_1449', 5356 bytes spill stores, 5336 bytes spill loads

638/640 0s 5ms/step -
accuracy: 0.8441 - loss: 0.4435

2025-04-28 12:16:14.032823: I
external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas
warning : Registers are spilled to local memory in function
'gemm_fusion_dot_1449_0', 176 bytes spill stores, 524 bytes spill loads

2025-04-28 12:16:14.055351: I
external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas
warning : Registers are spilled to local memory in function
'gemm_fusion_dot_1449', 12 bytes spill stores, 12 bytes spill loads

2025-04-28 12:16:14.177429: I
external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas
warning : Registers are spilled to local memory in function
'gemm_fusion_dot_1449', 104 bytes spill stores, 104 bytes spill loads

2025-04-28 12:16:14.418061: I
external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas
warning : Registers are spilled to local memory in function
'gemm_fusion_dot_1449', 5144 bytes spill stores, 5204 bytes spill loads

2025-04-28 12:16:14.671121: I
external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas
warning : Registers are spilled to local memory in function
'gemm_fusion_dot_1449', 4984 bytes spill stores, 4984 bytes spill loads

640/640 16s 16ms/step -
 accuracy: 0.8442 - loss: 0.4430 - val_accuracy: 0.9049 - val_loss: 0.2524
 Epoch 2/5
 640/640 4s 6ms/step -
 accuracy: 0.9194 - loss: 0.2184 - val_accuracy: 0.9076 - val_loss: 0.2571
 Epoch 3/5
 640/640 4s 6ms/step -
 accuracy: 0.9373 - loss: 0.1712 - val_accuracy: 0.9112 - val_loss: 0.2514
 Epoch 4/5
 640/640 4s 6ms/step -
 accuracy: 0.9513 - loss: 0.1369 - val_accuracy: 0.9119 - val_loss: 0.2494
 Epoch 5/5
 640/640 4s 6ms/step -
 accuracy: 0.9611 - loss: 0.1100 - val_accuracy: 0.9104 - val_loss: 0.2672

Fine-tuning: Descongelando desde la capa 100

- Capa 100 (block_11_expand_relu): Trainable = True
- Capa 110 (block_12_depthwise): Trainable = True
- Capa 120 (block_13_depthwise): Trainable = True
- Capa 130 (block_14_depthwise_relu): Trainable = True
- Capa 140 (block_15_project): Trainable = True
- Capa 149 (block_16_project): Trainable = True
- Capa 150 (block_16_project_BN): Trainable = True
- Capa 151 (Conv_1): Trainable = True
- Capa 152 (Conv_1_bn): Trainable = True
- Capa 153 (out_relu): Trainable = True

Re-compilando modelo para fine-tuning con LR bajo...

Iniciando fase de fine-tuning..

Epoch 1/12

640/640 25s 22ms/step -
 accuracy: 0.7880 - loss: 0.6939 - val_accuracy: 0.9273 - val_loss: 0.2613 -
 learning_rate: 1.0000e-05

Epoch 2/12

638/640 0s 7ms/step -
 accuracy: 0.9258 - loss: 0.2186
 Epoch 2: ReduceLROnPlateau reducing learning rate to 2.9999999242136253e-06.

640/640 6s 9ms/step -
 accuracy: 0.9258 - loss: 0.2186 - val_accuracy: 0.9127 - val_loss: 0.2671 -
 learning_rate: 1.0000e-05

Epoch 3/12

633/640 0s 7ms/step -
 accuracy: 0.9657 - loss: 0.1162
 Epoch 3: ReduceLROnPlateau reducing learning rate to 8.999999636216671e-07.

640/640 5s 8ms/step -
 accuracy: 0.9657 - loss: 0.1162 - val_accuracy: 0.9037 - val_loss: 0.2763 -
 learning_rate: 3.0000e-06

10 8. Modelo 5 — Data Augmentation

```
[20]: data_augmentation_layers = tf.keras.Sequential(
    [
        layers.RandomFlip("horizontal", seed=SEED),
        layers.RandomRotation(0.2, seed=SEED),
        layers.RandomZoom(0.2, seed=SEED),
        layers.RandomTranslation(height_factor=0.2, width_factor=0.2,
↪seed=SEED),
    ],
    name="data_augmentation",
)

try:
    base_model_for_aug = tl_model
except NameError:
    print("Cargando modelo base para aumentación desde archivo...")
    base_model_for_aug = tf.keras.models.load_model(MODELS_DIR / 'fine_tuning.
↪keras')

inputs = tf.keras.Input(shape=IMG_SIZE + (3,))
x = data_augmentation_layers(inputs, training=True)
outputs = base_model_for_aug(x)
aug_model = tf.keras.Model(inputs, outputs, name='model_with_augmentation')

aug_model.compile(optimizer=get_optimizer(1e-5),
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])

history_aug = aug_model.fit(
    train_ds,
    epochs=EPOCHS,
    validation_data=val_ds,
    callbacks=common_callbacks('data_aug')
)

json.dump(history_aug.history, open(HIST_DIR / 'data_aug.json', 'w'))
aug_model.save(MODELS_DIR / 'data_aug.keras')
```

Epoch 1/12

640/640 43s 57ms/step -

accuracy: 0.7094 - loss: 0.9075 - val_accuracy: 0.9076 - val_loss: 0.2659 -
learning_rate: 1.0000e-05

Epoch 2/12

640/640 34s 53ms/step -

accuracy: 0.7699 - loss: 0.6226 - val_accuracy: 0.9096 - val_loss: 0.2606 -
learning_rate: 1.0000e-05

Epoch 3/12

```

640/640          35s 55ms/step -
accuracy: 0.8076 - loss: 0.5249 - val_accuracy: 0.9064 - val_loss: 0.2574 -
learning_rate: 1.0000e-05
Epoch 4/12
640/640          0s 52ms/step -
accuracy: 0.8167 - loss: 0.4880
Epoch 4: ReduceLROnPlateau reducing learning rate to 2.9999999242136253e-06.
640/640          35s 54ms/step -
accuracy: 0.8167 - loss: 0.4880 - val_accuracy: 0.9041 - val_loss: 0.2669 -
learning_rate: 1.0000e-05
Epoch 5/12
639/640          0s 53ms/step -
accuracy: 0.8335 - loss: 0.4508
Epoch 5: ReduceLROnPlateau reducing learning rate to 8.999999636216671e-07.
640/640          35s 55ms/step -
accuracy: 0.8336 - loss: 0.4508 - val_accuracy: 0.9053 - val_loss: 0.2604 -
learning_rate: 3.0000e-06

```

11 9. Evaluación en test y comparativa final

```

[23]: def evaluate_and_log(model_path, name):
        model = tf.keras.models.load_model(model_path)
        loss, acc = model.evaluate(test_ds, verbose=0)
        return {'Model_Name': name, 'Test_Accuracy': acc}

[25]: results = pd.DataFrame([
    evaluate_and_log(MODELS_DIR/'base_model.keras', 'base_model'),
    evaluate_and_log(MODELS_DIR/'advance_model.keras', 'advance_model'),
    evaluate_and_log(MODELS_DIR/'hp.keras', 'hp'),
    evaluate_and_log(MODELS_DIR/'fine_tuning.keras', 'fine_tuning'),
    evaluate_and_log(MODELS_DIR/'data_aug.keras', 'data_aug')
])


display(results.style.background_gradient(cmap='Reds', subset=['Test_Accuracy'])
        .format({'Test_Accuracy': '{:.3f}'}))

plt.figure(figsize=(10,5))
for m, color in zip(['base_model', 'advance_model', 'hp', 'fine_tuning', 'data_aug'],
                    ['tab:blue', 'tab:orange', 'tab:green', 'tab:red', 'tab:
purple']):
    hist = json.load(open(HIST_DIR/f'{m}.json'))
    plt.plot(hist['val_accuracy'], label=m.replace('_', ' ').title(),
             linewidth=2)
plt.legend(); plt.ylabel('Validation accuracy'); plt.xlabel('Epoch'); plt.
ylim(0,1); plt.grid(alpha=.3)

```

```
plt.tight_layout(); plt.show()
```

<pandas.io.formats.style.Styler at 0x7f2d9dc26b40>



12 10. Visualización rápida de predicciones con matplotlib

```
[27]: model = tf.keras.models.load_model(MODELS_DIR / 'data_aug.keras')
model.summary(line_length=80)

NUM_IMAGES = 12
test_iter = test_ds.unbatch().take(NUM_IMAGES)

plt.figure(figsize=(12, 9))
for idx, (img, true_lab) in enumerate(test_iter):
    pred_prob = model.predict(img[tf.newaxis, ...], verbose=0)
    pred_label = tf.argmax(pred_prob, axis=1).numpy()[0]

    ax = plt.subplot(3, 4, idx + 1)
    plt.imshow(img.numpy().astype("uint8"))
    ax.axis("off")

    correct = (pred_label == true_lab.numpy())
    color = "green" if correct else "red"
    ax.set_title(
        f"GT: {class_names[true_lab]} \nPred: {class_names[pred_label]}",
        fontsize=9, color=color, pad=4
    )
```

```
plt.tight_layout()
plt.show()
```

Model: "model_with_augmentation"

Layer (type)	Output Shape	Param #
input_layer_4 (InputLayer)	(None , 150, 150, 3)	0
data_augmentation (Sequential)	(None , 150, 150, 3)	0
MobileNetV2_finetune (Functional)	(None , 6)	2,422,726

Total params: 6,475,096 (24.70 MB)


Trainable params: 2,026,182 (7.73 MB)

Non-trainable params: 396,544 (1.51 MB)

Optimizer params: 4,052,370 (15.46 MB)

2025-04-28 12:23:39.424899: I tensorflow/core/framework/local_rendezvous.cc:407]
Local rendezvous is aborting with status: OUT_OF_RANGE: End of sequence


GT: sea
Pred: sea



GT: street
Pred: street



GT: forest
Pred: forest



GT: buildings
Pred: buildings



GT: sea
Pred: sea



GT: glacier
Pred: glacier



GT: mountain
Pred: mountain



GT: glacier
Pred: glacier




GT: glacier
Pred: glacier



GT: street
Pred: street



GT: glacier
Pred: glacier



GT: forest
Pred: forest

